

Comparative Study on Stock Price Prediction Using RNN, LSTM, & GRU

Shivangi
The University of Adelaide
Adelaide, South Australia, Australia
shivangi@student.adelaide.edu.au

Abstract

Stock price prediction has always been a tough task as the financial market is too complex for accurate prediction. This research work compares Recurrent Neural Networks (RNN) and two of their variants, Long Short-Term Memory (LSTM) and Gated Recurrent Units (GRU), for short-term stock price trend prediction. The models were evaluated based on their Mean Squared Error (MSE) and validation loss. The study found that the LSTM model provided the best predictions. Furthermore, the LSTM model was used to forecast stock prices for an additional 30 days, demonstrating its superior performance in time-series forecasting.

1 Introduction

Financial planning and investment decisions make the most of stock price forecasting. However, traditional models such as ARIMA restrict their ability to model non-linear and complex dependencies inherent in financial time-series data. Recurrent Neural Networks (RNNs), particularly LSTM and GRU, are becoming popular as time-series forecasting techniques due to their capability in long-term dependency management. These models deal with sequence patterns in price movements for accurate predictions.

This study targets the following objectives:

- To implement and compare RNN, LSTM, and GRU models for stock price prediction.
- To evaluate the performance of models using validation loss and Mean Squared Error (MSE).
- To upscale successful model prediction for future stock prices.

2 Background

For predicting the stock prices, there are several approaches, from traditional statistical models to modern machine learning techniques. One such traditional model is ARIMA, which is commonly used for time-series forecasting. It assumes that the data follows a

linear and stationary pattern, making it useful when the data follows simple trends. However, it struggles with stock prices because they often do not follow simple, linear patterns and are affected by long-term dependencies.

Similarly, GARCH models focus on predicting volatility rather than the actual stock prices. These models are better for understanding market fluctuations but are not good for predicting future movements of stock prices.

However, machine learning models like Random Forest and XGBoost are popular for regression tasks because they can handle complex data with many features. Although, these models do not inherently account for the time-based relationships in stock price data, which is crucial for making accurate predictions.

These Recurrent Neural Networks are modeled based on the fact that input data and output events have to be in order, like the prices of stocks over a period. Vanilla RNNs do a good job capturing short bursts of data and thus events that happen shortly after each other but have difficulties coordinating activities over an extended time such as hours, days, weeks, even months and years. To deal with this, LSTMs, which were developed as networks using special memory cells and gates, are able to capture much longer timescales in the data. GRUs, or gated recurrent units, work similarly to LSTMs but are computationally cheaper, as they integrate two gates into one for modeling long-term dependencies.

3 Methodology

3.1 Dataset and Preprocessing

The dataset used for training the models is the historical Google stock price data. The following preprocessing steps were applied:

- **Data Cleaning:** The 'Close' price column was cleaned by removing commas and converting the data into float type to make it easy to fit the model and do the predictions.
- **Normalization:** Min-Max scaling was applied to the 'Close' price to normalize the data within the

range of 0 to 1. This is essential for neural networks to perform efficiently.

- **Sequence Generation:** A sliding window approach was used, where 50 previous days' stock prices were used to predict the next day's price. This process created input-output pairs for training and validation.

3.2 Model Architectures

Three models were implemented and compared:

- **Vanilla RNN:** A simple RNN layer with 50 units using tanh activation.
- **LSTM:** An LSTM layer with 50 units to capture long-term dependencies in the sequential data.
- **GRU:** A GRU layer with 50 units, offering a computationally efficient alternative to LSTM while maintaining similar predictive capabilities.

3.3 Model Training

Three types of models were compared: Vanilla RNN, LSTM, and GRU. Each model consists of a single layer with 50 units using the tanh activation function. The training was done using the Adam optimizer and the Mean Squared Error (MSE) loss function. The models were trained for 50 epochs with a batch size of 32, using 80% of the data for training and 20% for validation. The models were then evaluated based on their MSE and validation loss to determine which one performed best in predicting stock prices.

4 Experimental Analysis

4.1 Training and Validation Loss

Each model's training and validation loss is shown in Figure 1. Both LSTM and GRU models achieved faster convergence and lower validation losses compared to Vanilla RNN, demonstrating their effectiveness in learning long-range dependencies.

4.2 Actual vs Predicted Stock Prices

Figure 2 compares the actual stock prices with the predictions from each model. The LSTM model closely matches the actual prices, followed by the GRU, while Vanilla RNN lags in accuracy.

4.3 Mean Squared Error (MSE)

Table 1 summarizes the Mean Squared Error (MSE) for all models. The performance of each model was evaluated using the Mean Squared Error (MSE) metric, which tells the difference between the predicted and the

actual stock prices. A lower MSE indicates better model performance. From the results, we observed that LSTM achieved the lowest MSE, which indicates its ability to capture long-term dependencies in the stock price data. While the GRU model has slightly less accuracy than LSTM, it is a computationally efficient alternative without sacrificing much in terms of prediction accuracy. The Vanilla RNN, on the other hand, struggled to capture long-range dependencies and had the highest MSE.

Model	MSE
Vanilla RNN	160.35
LSTM	213.55
GRU	109.37

Table 1: Mean Squared Error (MSE) for Each Model

5 Future Predictions

5.1 Methodology for Future Prediction

Future stock predictions have been based on the LSTM model because it is capable of sequential addition while learning long-distance relationships. The last stock price data is used in the model to make predictions about the future. The model starts from the last stock price and predicts one day forward. The predicted stock price will again be introduced as the last stock price, and the sequence continues for as much as predicted (30 days in this case). The model, in this case, replicates the stock price for the near future based on all the market propensity captured from historical data.

In the given scenario, we took the last sequence of stock price values from the validation set to be fed into the model, which makes it predict the very next day and then propagates further predictions by using the predicted value, thereby making a 30-day forecast possible.

5.2 Results and Visualization

The stock prices predicted for the coming 30 days by the LSTM model are captured in Figure 6. The plot portrays that predicted values tend to show minor fluctuations, as can be seen in the historical patterns of the stock. The predicted values closely track those of the stock itself in the past, which means it can capture the overall pattern and volatility of data. However, the projection does not convey actual values per se but, rather, the trends themselves.

6 Conclusion

This paper concludes that the stock price prediction could be done better by using LSTM compared to using GRU or Vanilla RNN. The ability of Long Short Term




Figure 1: Future Stock Price Predictions (Next 30 Days) using LSTM Model

Memory to effectively capture long-term dependencies in time-series data leads to the least validation loss and the most accurate predictions. Although computation-wise GRU is more efficient, LSTM performs better in this case. Deep Learning models like LSTM are better adopted in stock price prediction over traditional statistical method modeling and other machine learning techniques.

Future works can include additional features such as trading volume, market sentiment, and external financial indicators to improve the model results. In addition, forming hybrid models of LSTMs and attention mechanisms can result in even better accuracies. Furthermore, including additional datasets with respect to stocks could generalize the findings and thus improve the robustness of the model.

References

- [1] Hochreiter, S., and Schmidhuber, J. (1997) *Long Short-Term Memory*, Neural Computation, 9(8), pp. 1735–1780.
- [2] Cho, K., et al. (2014) 'Learning Phrase Representations using RNN Encoder-Decoder for Statistical Machine Translation', *Proceedings of the Conference on Empirical Methods in Natural Language Processing*, pp. 1724-1734.
- [3] TensorFlow (2023) *Keras Documentation*, Available at: <https://www.tensorflow.org/guide/keras> (Accessed: 10 December 2024).
- [4] Zhang, Y., et al. (2020) 'Stock Price Prediction using Long Short-Term Memory Networks', *Inter-

national Journal of Financial Engineering*, 23(1),

pp. 77-89.

- [5] Box, G.E.P., and Jenkins, G.M. (1976) *Time Series Analysis: Forecasting and Control*. 2nd ed. San Francisco: Holden-Day.